

PDE-Refiner: Achieving Accurate Long Rollouts with Neural PDE Solvers

Phillip Lippe^{1,2}, Bas S. Veeling¹, Paris Perdikaris¹, Richard E. Turner¹, Johannes Brandstetter¹

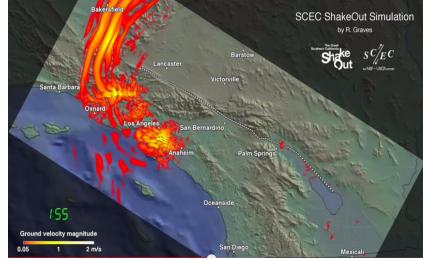
¹Microsoft Research AI4Science, ²University of Amsterdam



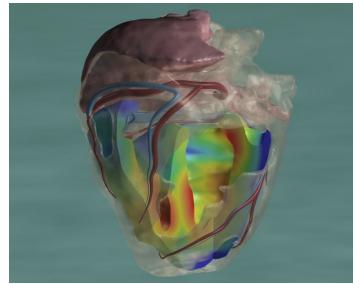
Project Website



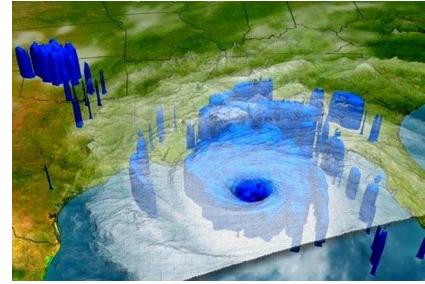
(Large-scale) PDE systems are ubiquitous



Earthquakes



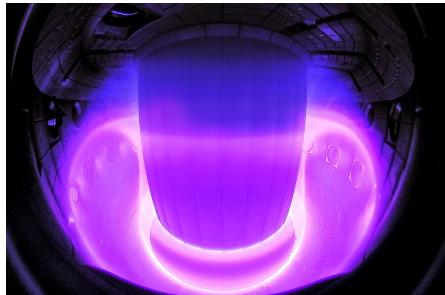
Heart dynamics



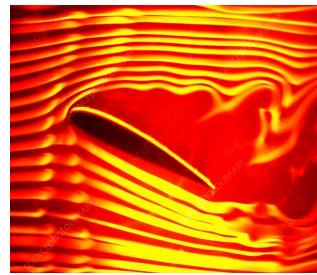
Weather prediction



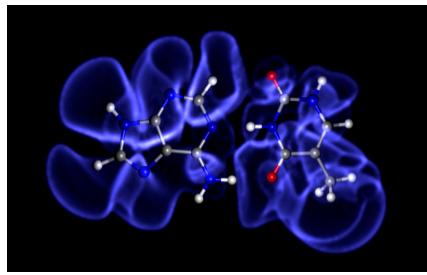
Galaxy collisions



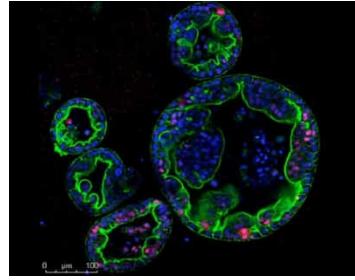
Plasma physics



Airplane design



Electronic structure



Tumor growth

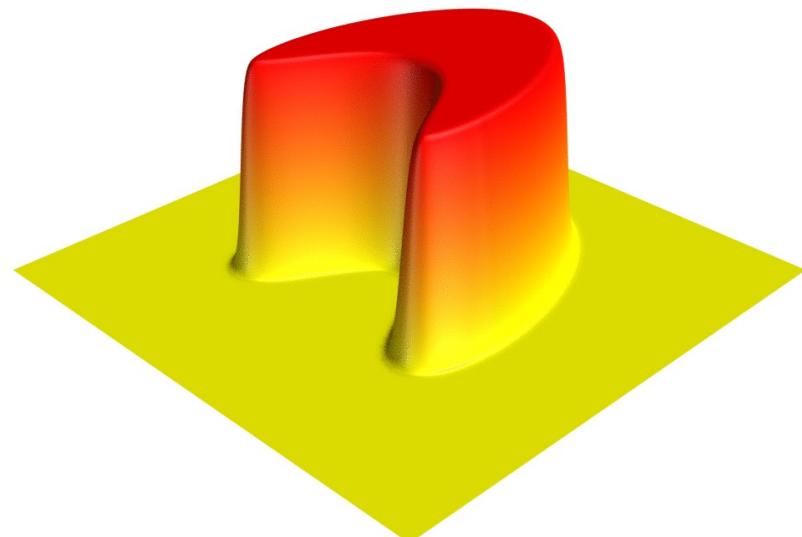
Solving PDEs – the traditional way

- Formulation of time dependent Partial Differential Equations (PDEs):

$$\partial_t \mathbf{u} = F(t, \mathbf{x}, \mathbf{u}, \partial_{\mathbf{x}} \mathbf{u}, \partial_{\mathbf{xx}} \mathbf{u}, \dots) \quad (t, \mathbf{x}) \in [0, T] \times \mathbb{X}$$

$$\mathbf{u}(0, \mathbf{x}) = \mathbf{u}^0(\mathbf{x}), \quad B[\mathbf{u}](t, x) = 0 \quad \mathbf{x} \in \mathbb{X}, (t, \mathbf{x}) \in [0, T] \times \partial \mathbb{X}$$

- Partition spatial and temporal domain into grid
- Estimate spatial derivatives, e.g., via finite difference
- Solve time derivative with classical ODE solvers,
e.g., Runge-Kutta methods



Example: Heat Equation ([credit](#))

Example PDEs

- 2D Kolmogorov Flow (KM)

- Fluid Dynamics
- Incompressible Navier-Stokes
- Known for its chaotic behavior
- Accurate solving requires expensive, high resolution

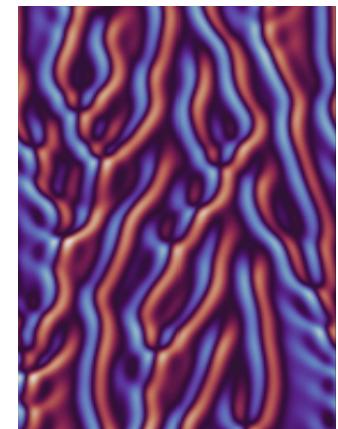
$$\partial_t \mathbf{u} + \nabla \cdot (\mathbf{u} \otimes \mathbf{u}) = \nu \nabla^2 \mathbf{u} - \frac{1}{\rho} \nabla p + \mathbf{f}$$



- 1D Kuramoto-Sivashinsky Equation (KS)

- Fourth-order nonlinear PDE
- Fluid Dynamics, e.g., plasmas and flame propagation
- Rich dynamical characteristics and chaotic behavior

$$u_t + uu_x + u_{xx} + \nu u_{xxxx} = 0$$



Challenges for classical solvers

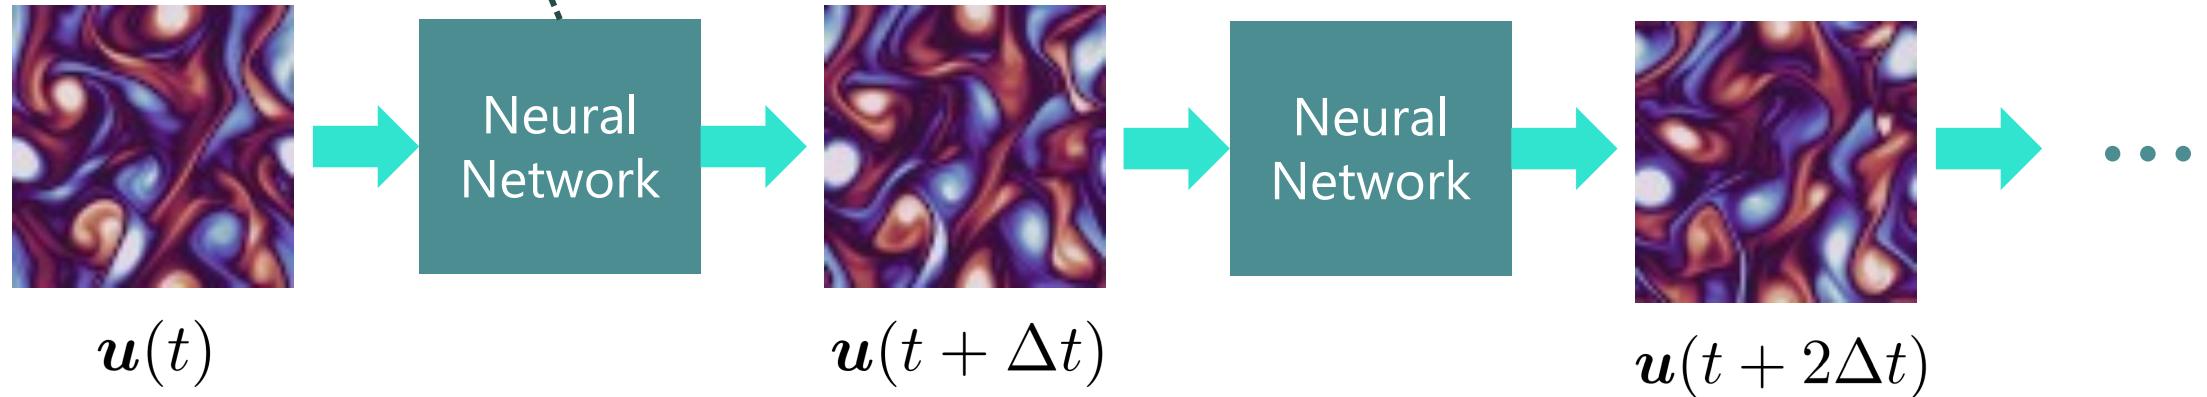
- Small errors have large long-term impact in chaotic PDEs
- Requires small time steps and/or large resolution \Rightarrow expensive
- Example: Kolmogorov Flow
 - Requires resolution of 2048x2048
 - Time step of 0.007s
 - 20 second rollout takes >30 minutes on an A100
- Can we use ML to solve PDEs more efficiently?



Neural PDE Solvers

- Neural Operators learn to predict future solutions

$$\mathbf{u}(t + \Delta t) = \mathcal{G}_t(\Delta t, \mathbf{u}(t))$$



- Trained on one-step predictions
- Long horizon predictions via autoregressive rollout

Neural PDE Solvers - Desiderata

1. Long-Horizon Accuracy

- Remain close to ground truth solution for long time

2. Long-Horizon Stability

- Generate physically realistic solutions and not diverge

3. Uncertainty Estimation

- Know when not to trust your neural surrogate anymore

Neural PDE Solvers - Desiderata

1. Long-Horizon Accuracy

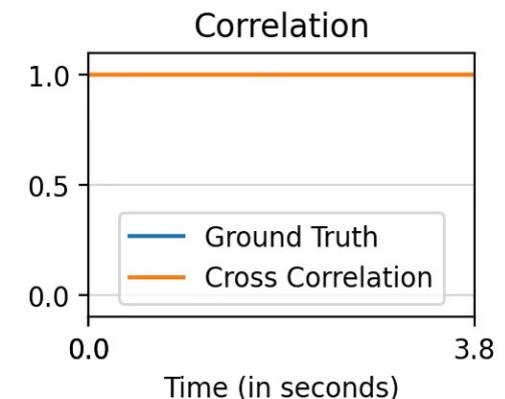
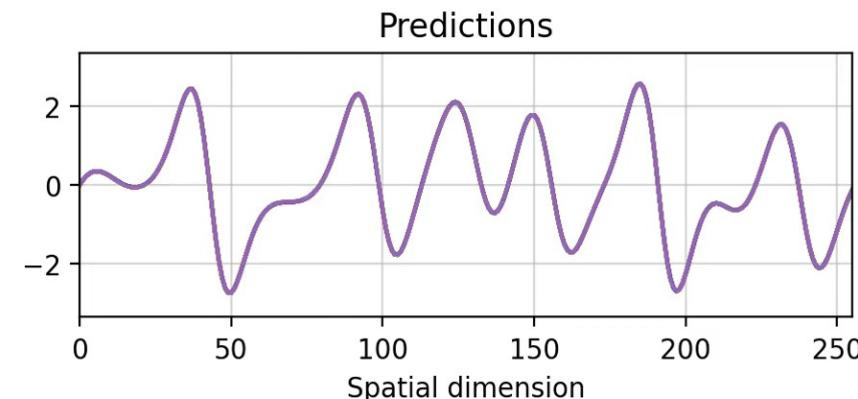
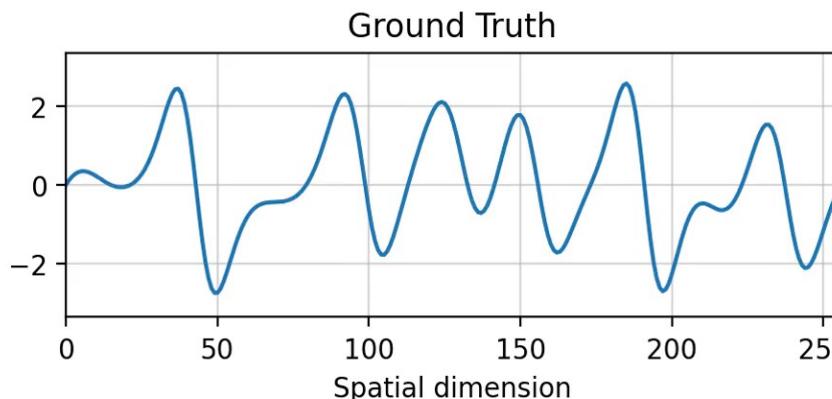
- Remain close to ground truth solution for long time

2. Long-Horizon Stability

- Generate physically realistic solutions and not diverge

3. Uncertainty Estimation

- Know when not to trust your neural surrogate anymore



Challenges of Accurate Long Rollouts

Training Neural PDE Solvers

- Commonly trained with one-step MSE:

$$\mathcal{L}_{\text{MSE}} = \|u(t) - \text{NO}(u(t - \Delta t))\|^2$$

- Tradeoff in time step size

- Large time steps give fast solvers, but harder to learn
 - Small time steps are easier to learn and generalize, but require many autoregressive steps
 - In practice, small time steps commonly achieve better performance

- Evaluate on rolling out model on its own predictions

- Check when it diverges from ground truth, e.g., in terms of correlation of MSE loss

Errors during Autoregressive Rollout

Three main sources of error during rollout:

1. Error Accumulation

- Models estimate the temporal difference between time steps, added to the original input
- Errors on the initial input are forwarded to future steps during rollout

2. Error Propagation

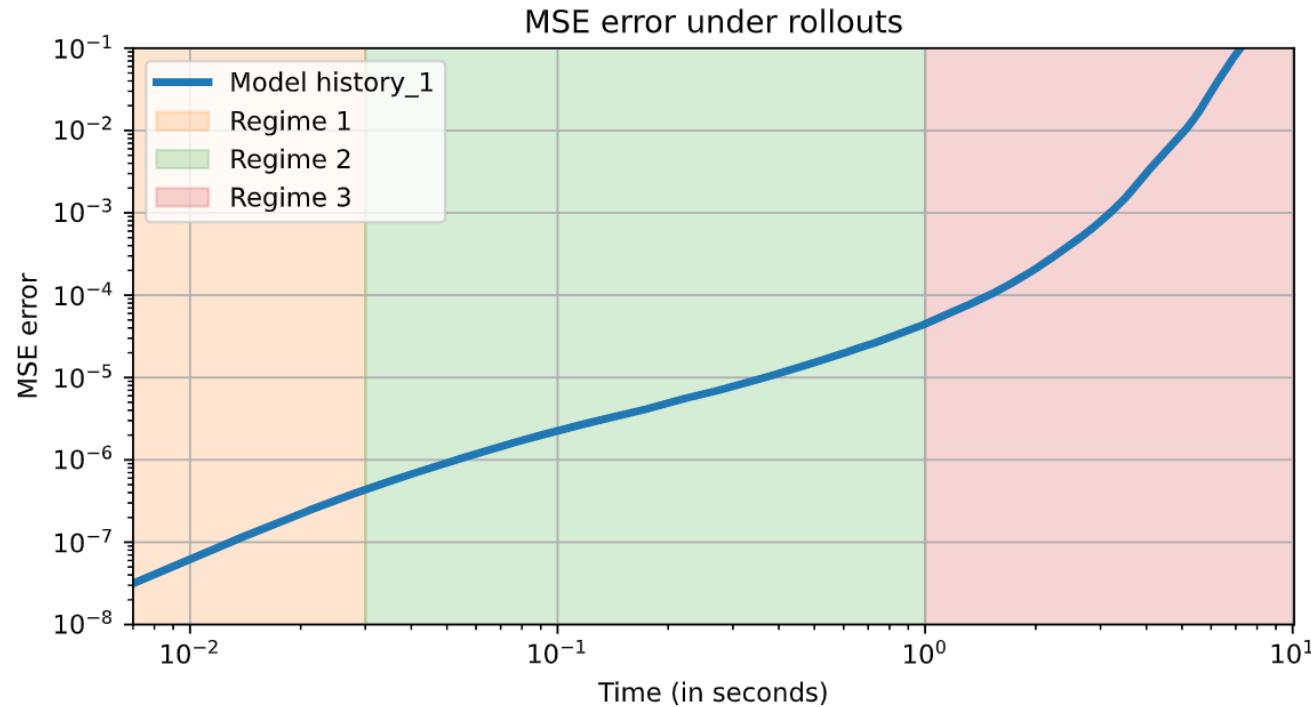
- Errors on network input influence the predicted temporal difference / dynamics

3. Input Distribution Shift

- Models are trained on ground truth data
- Predictions with errors may have a different distributions
- Can cause the model to diverge and make arbitrary predictions

Autoregressive Error Propagation

2D Kolmogorov Flow



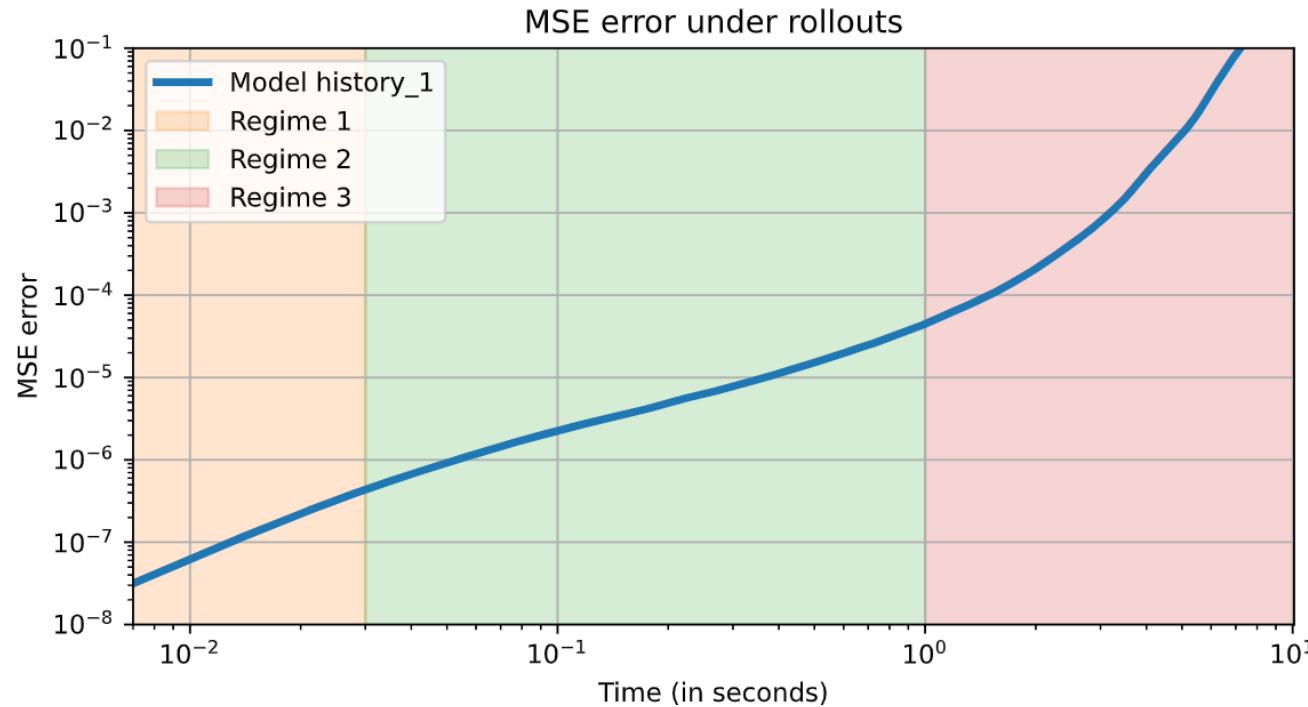
Regime 1

- Only error accumulation
- Consecutive errors are highly correlated \Rightarrow quadratic increase of error

$$\varepsilon_t \sim t^2 \varepsilon_0$$

Autoregressive Error Propagation

2D Kolmogorov Flow



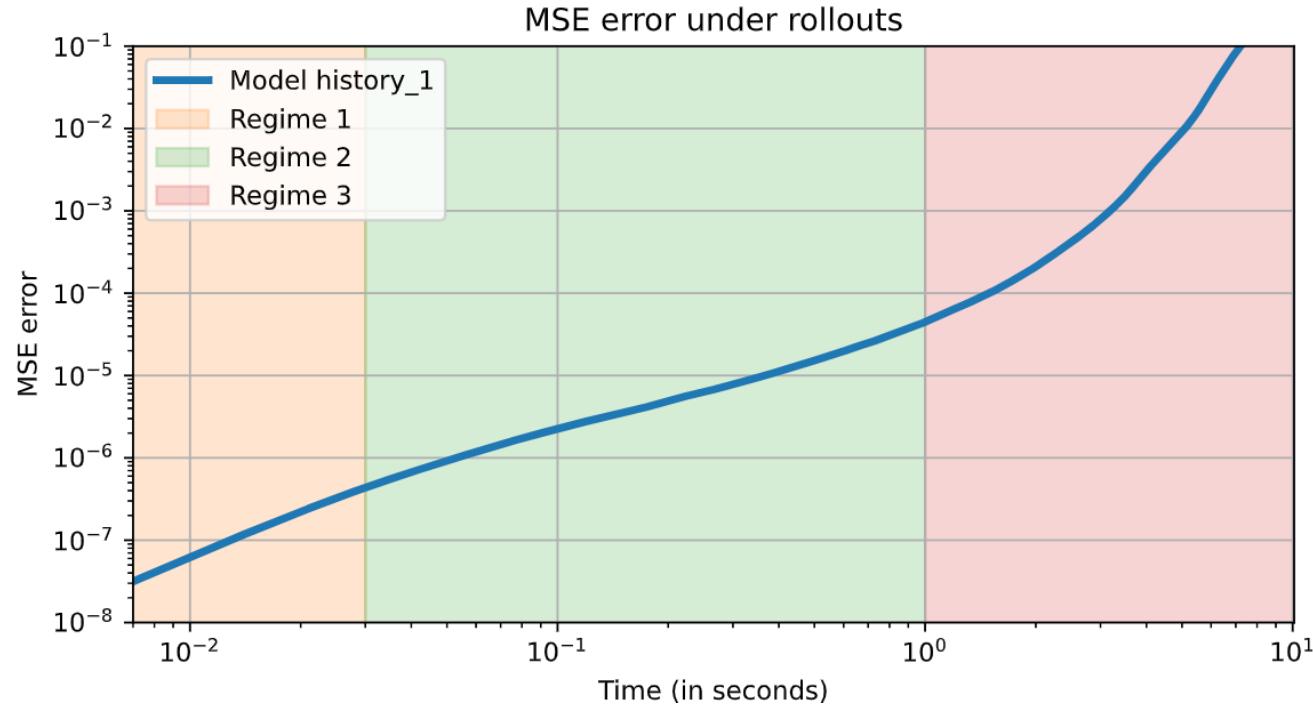
Regime 2

- Mainly error accumulation
- Errors become less correlated
⇒ slower increase of error

$$\varepsilon_t \sim t^{1.5} \varepsilon_0$$

Autoregressive Error Propagation

2D Kolmogorov Flow

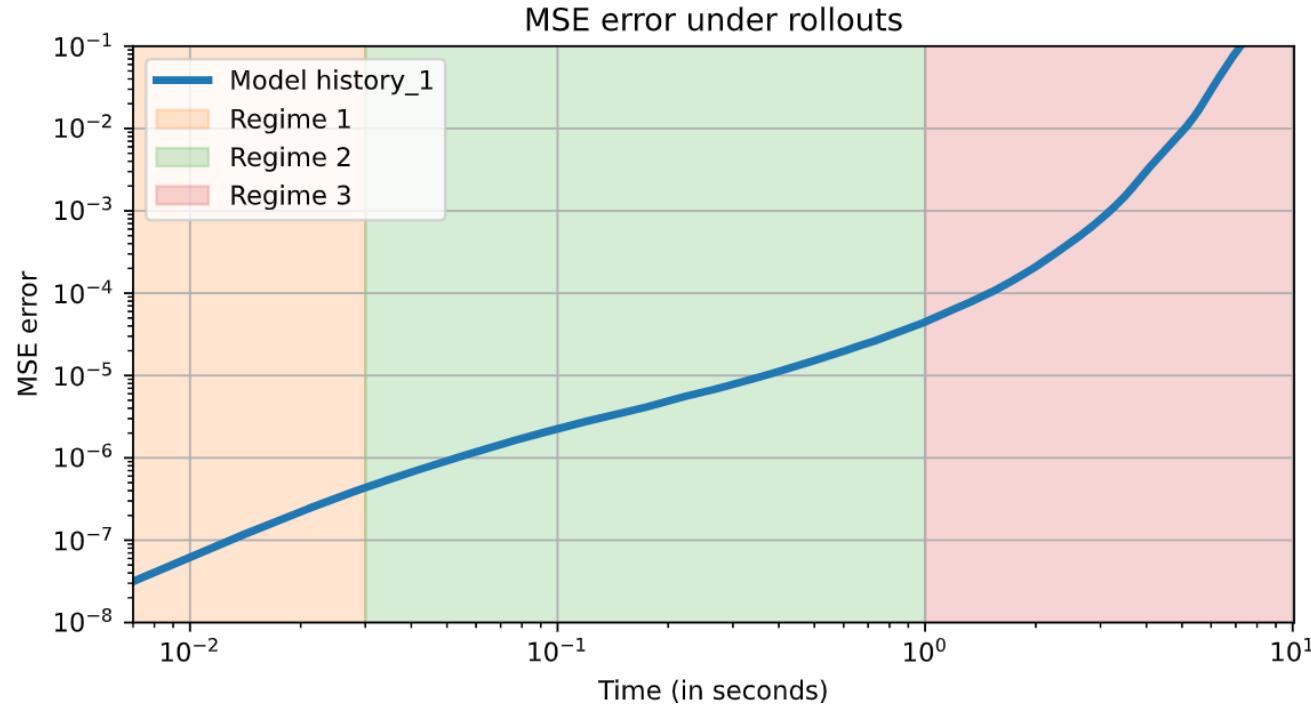


Regime 3

- Error propagation dominates
- Errors exponentially increase, diverge from ground truth
- Predictions yet appear physically realistic

Autoregressive Error Propagation

2D Kolmogorov Flow

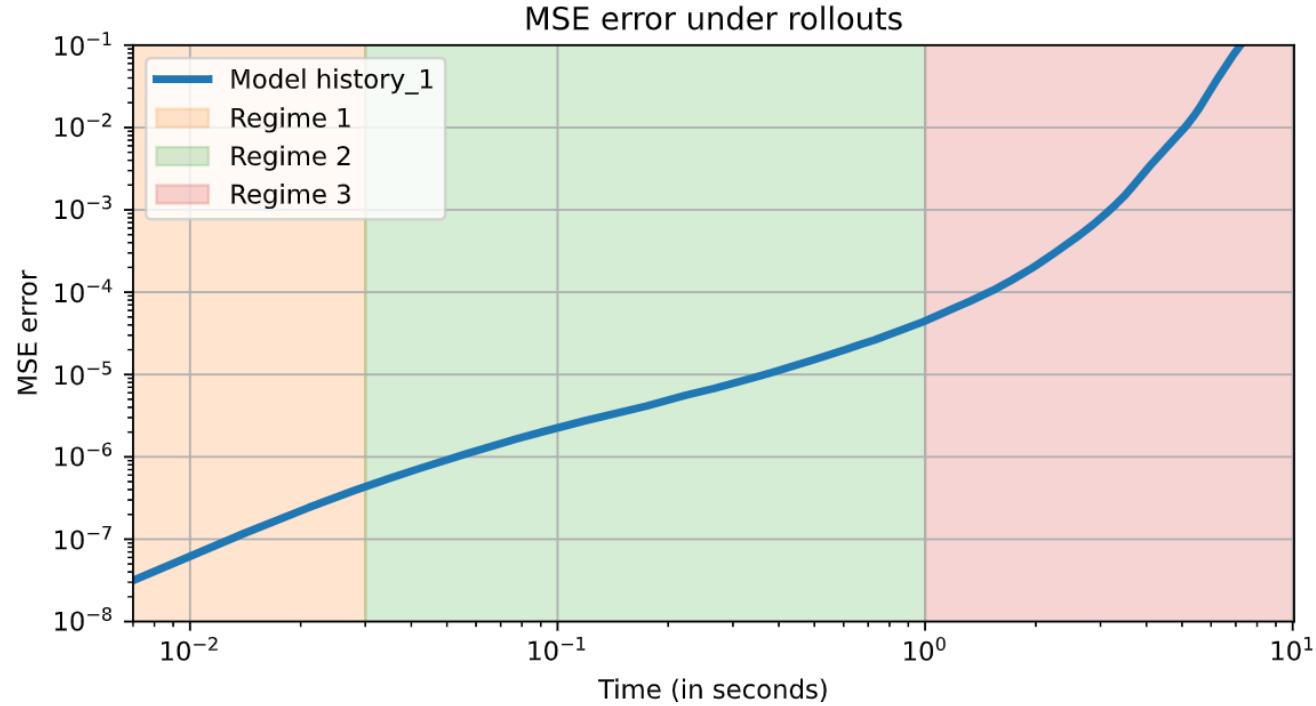


Regime 4

- Input Distribution Shift occurs much later in our setup
- Often after > 10 times longer rollout than divergence time

Autoregressive Error Propagation

2D Kolmogorov Flow



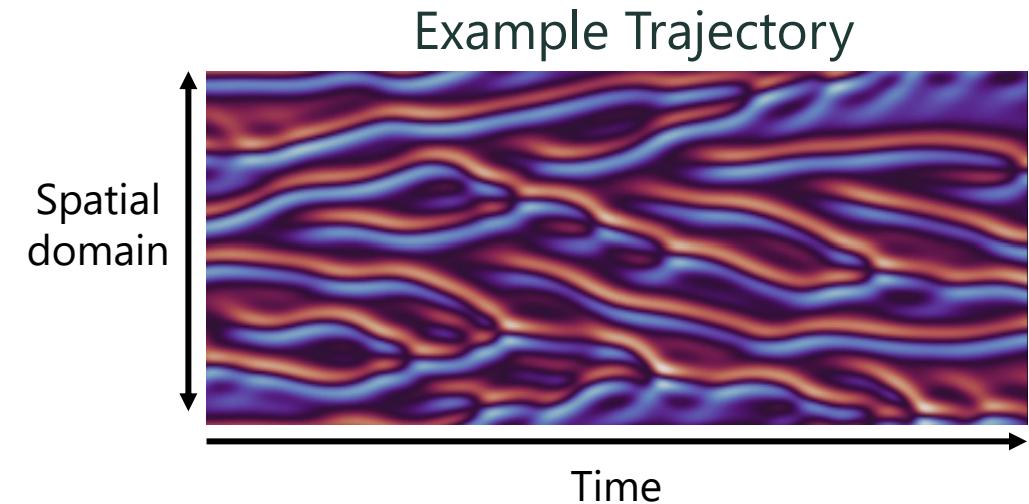
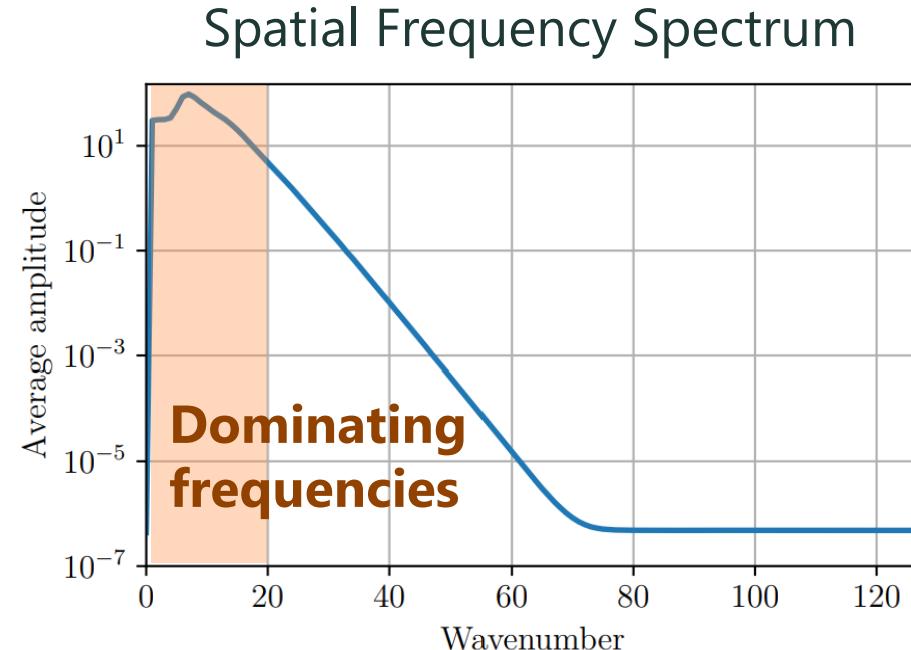
Rollout potentially improves by:

- Lower one-step loss
- Delaying error propagation (regime 3)

Delaying Error Propagation

- Example: 1D Kuramoto-Sivashinsky equation (KS)

$$u_t + uu_x + u_{xx} + \nu u_{xxxx} = 0$$



Delaying Error Propagation

- Example: 1D Kuramoto-Sivashinsky equation (KS)

$$u_t + \boxed{uu_x} + \boxed{u_{xx} + \nu u_{xxxx}} = 0$$

Non-linear term causes all spatial frequencies to interact long-term

High-order derivatives increase importance of high frequencies in spatial domain

→ For long accurate rollouts, model **all** spatial frequencies accurately
Errors in higher frequencies have low short-term, but **high long-term impact**

Delaying Error Propagation

- Many challenging PDEs follow similar pattern, for example:

KS equation $u_t + \boxed{uu_x} + \boxed{u_{xx} + \nu u_{xxxx}} = 0$

Burger's equation $u_t + \boxed{uu_x} + \boxed{u_{xx}} = 0$

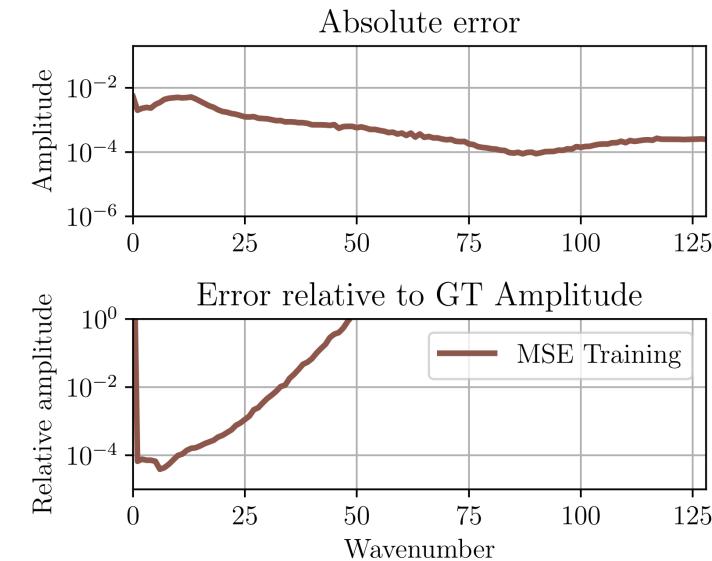
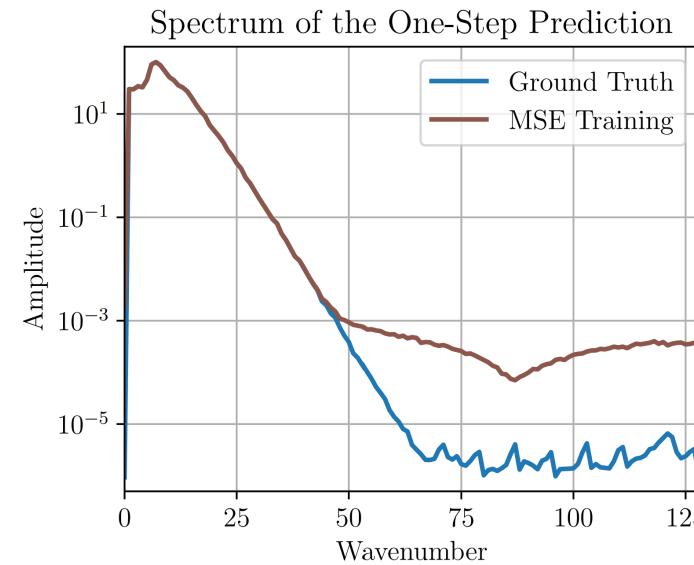
Korteweg-de Vries equation $u_t + \boxed{6uu_x} + \boxed{u_{xxx}} = 0$

KdV-Burger's equation $u_t + \boxed{2uu_x} - \boxed{\nu u_{xx} + \mu u_{xxx}} = 0$

→ For long accurate rollouts, model **all** spatial frequencies accurately
Errors in higher frequencies have low short-term, but **high long-term impact**

Case Study: Kuramoto-Sivashinsky

- How well do MSE-trained surrogates cover the frequency spectrum?



- Neural surrogates focus on **dominating** frequencies, losing high frequencies
- Inherently limits the maximum rollout time

Challenges of Accurate Long Rollouts – Summary

- Main causes for divergence: error accumulation and error propagation
- History information improves one-step, but accelerates error propagation
- MSE surrogates poorly model low-amplitude frequencies, inevitably setting a maximum possible rollout time



How can we capture the **whole** frequency spectrum better?

PDE-Refiner

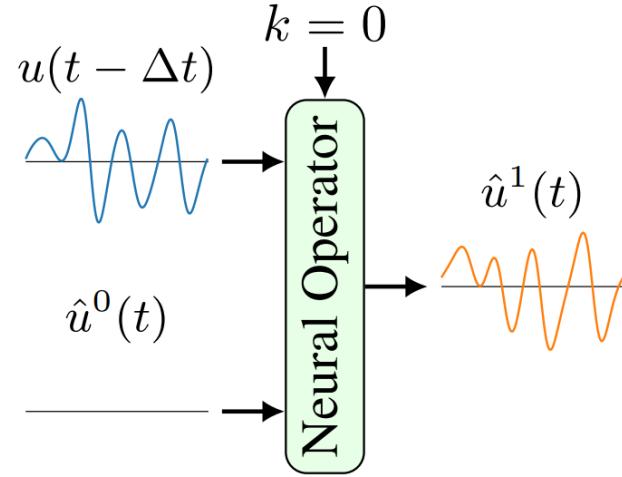
Achieving Accurate Long Rollouts via an Iterative Refinement Process

Idea

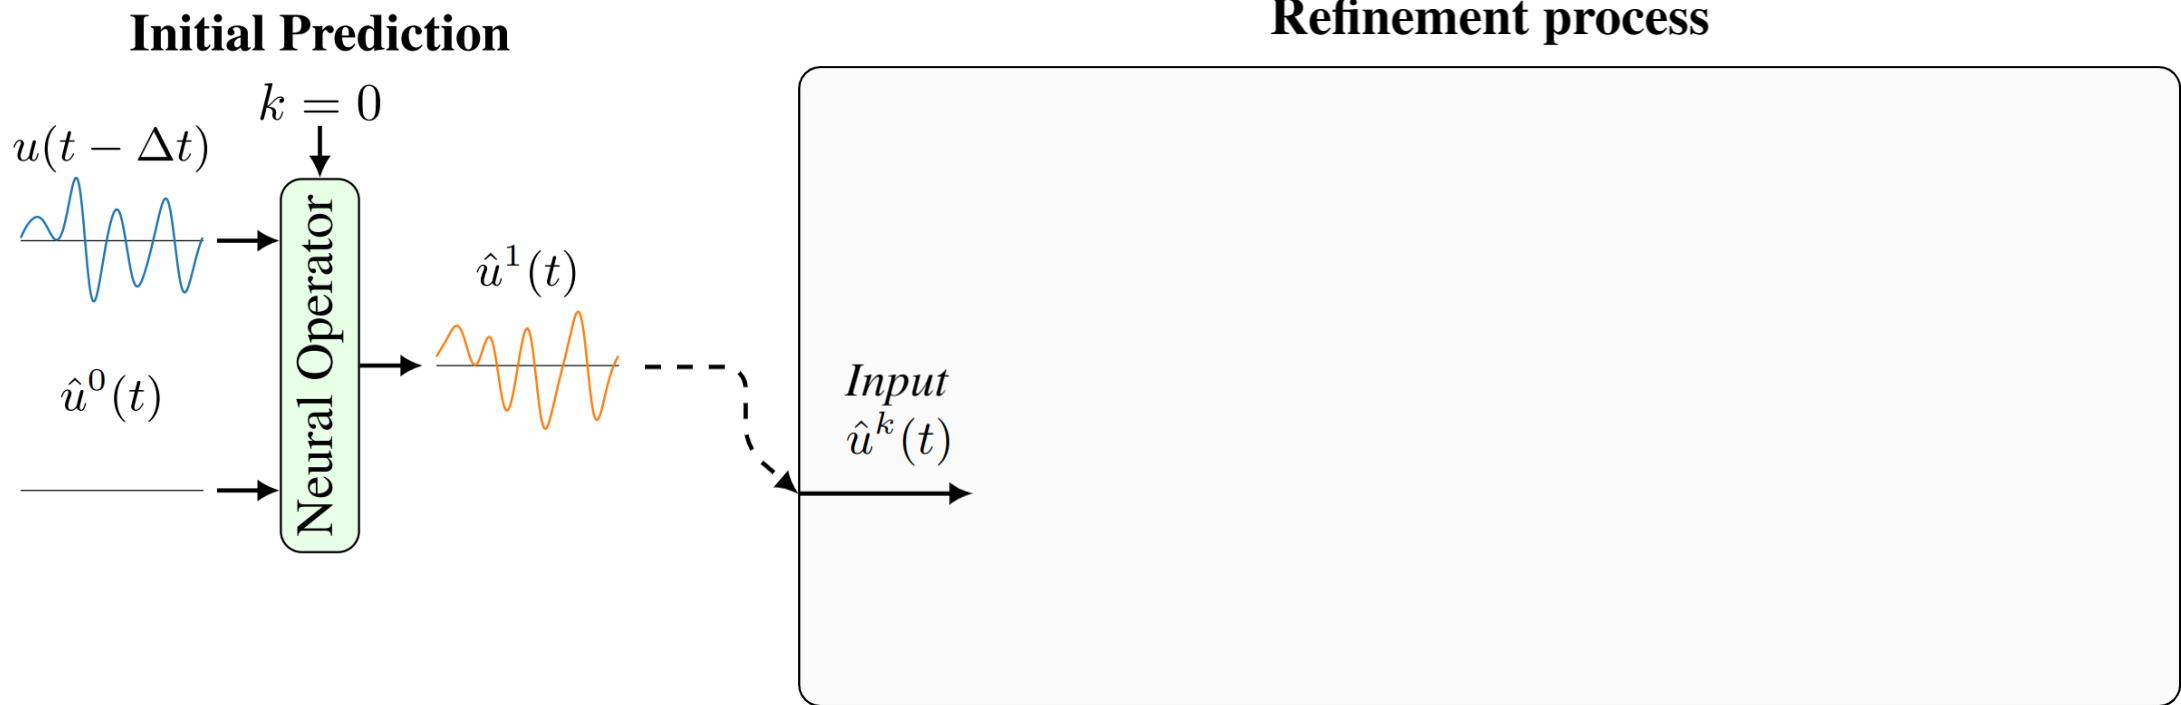
- Goal: improve prediction of low-amplitude frequencies
- Difficult to predict all frequencies perfectly at once
 - ⇒ iterative refinement process to finetune the prediction step-by-step
- At each refinement step, focus on information/error below a certain amplitude
 - Implemented via a denoising objective
- Use multiple refinement steps to cover larger spectrum

PDE-Refiner

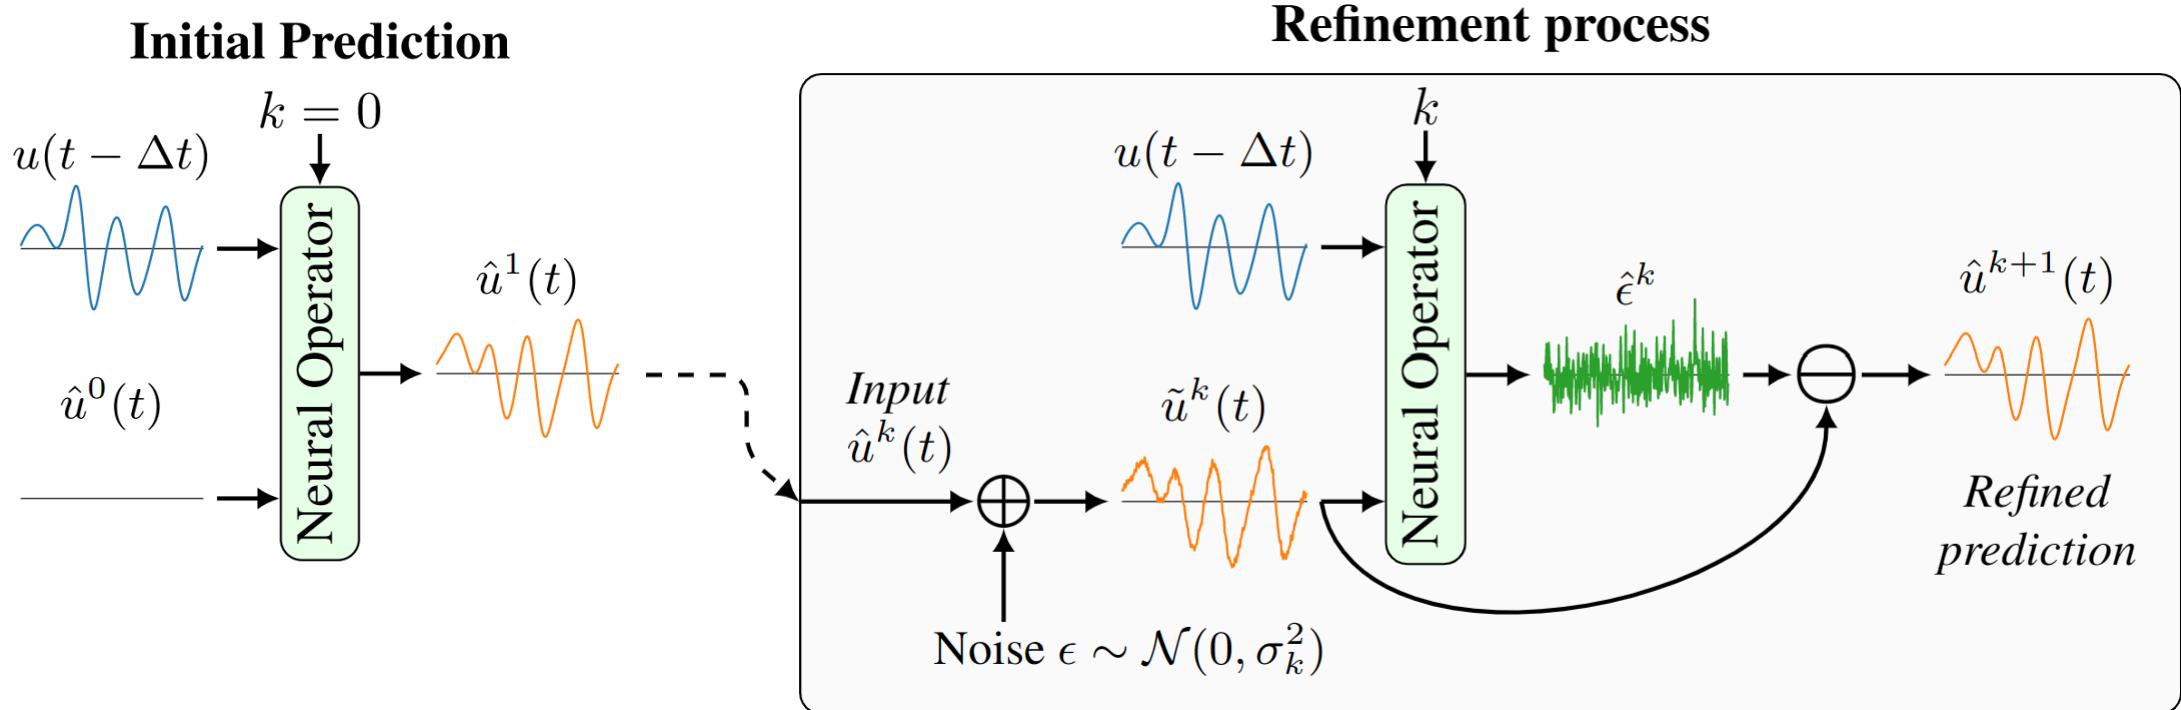
Initial Prediction



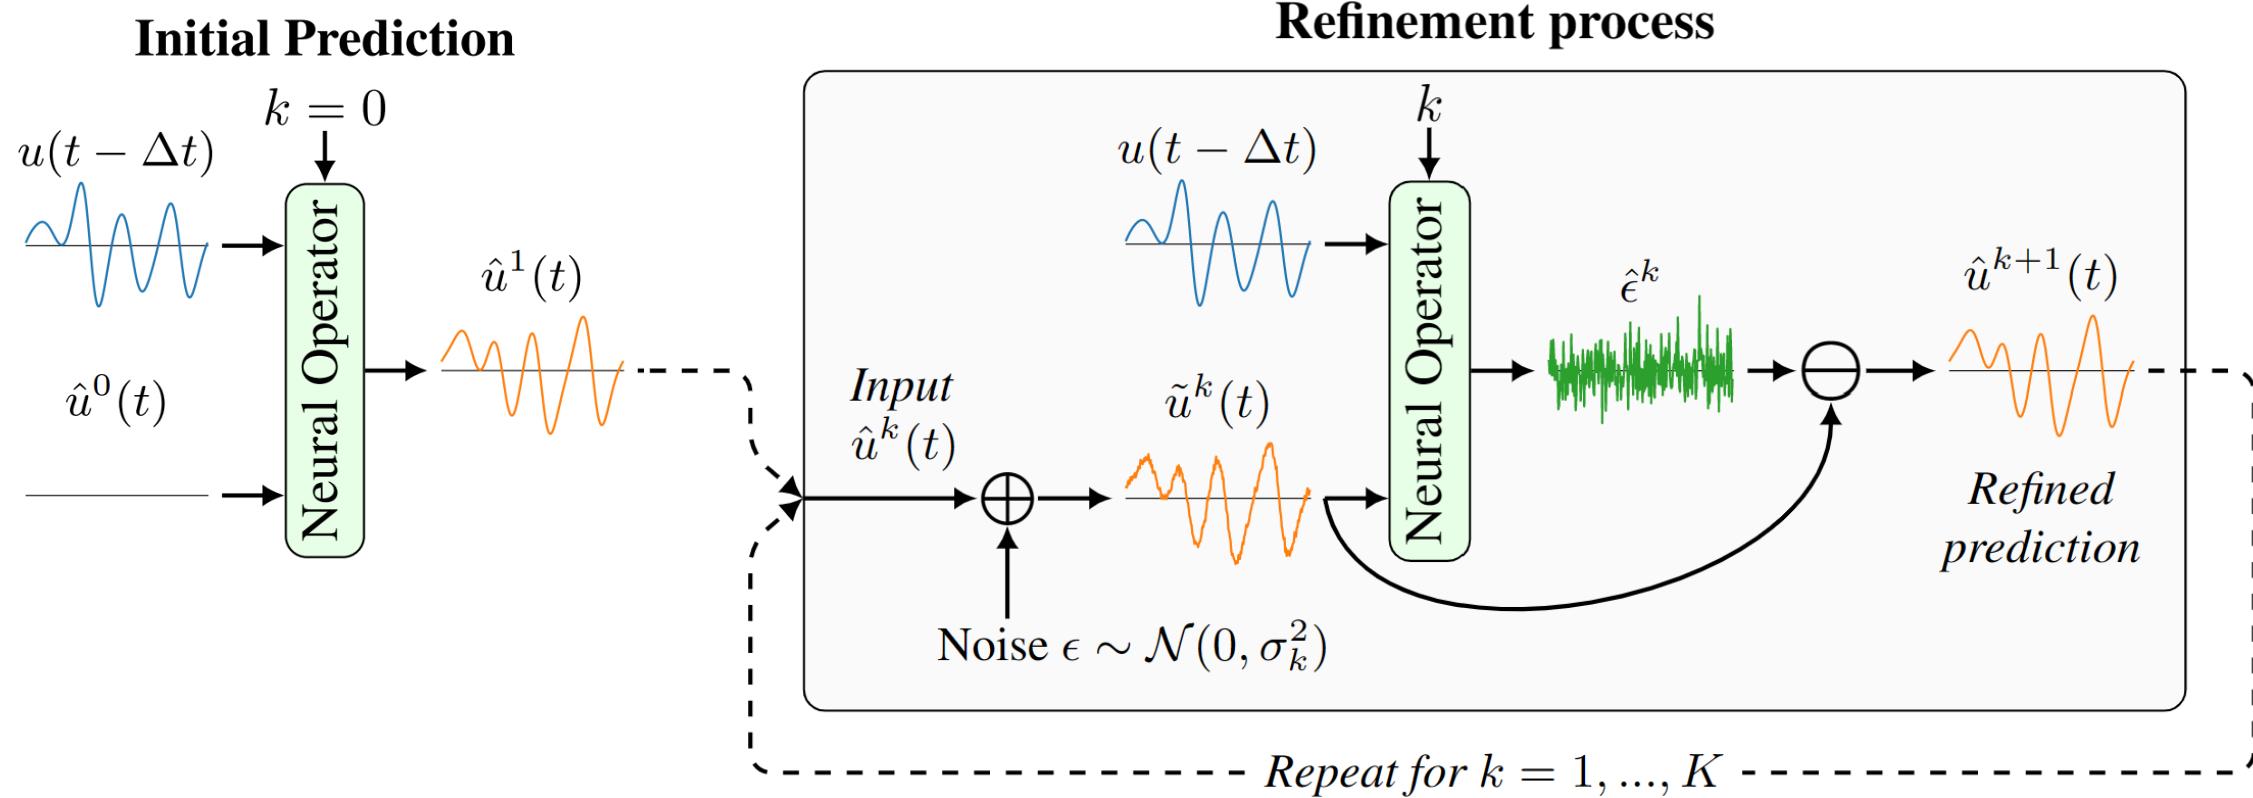
PDE-Refiner



PDE-Refiner



PDE-Refiner



PDE-Refiner – Training

- Initial prediction: common MSE objective

$$\mathcal{L}^0(u, t) = \|u(t) - \text{NO}(\hat{u}^0(t), u(t - \Delta t), 0)\|_2^2$$

$\nearrow = 0$

- Refinement steps: denoise ground truth data

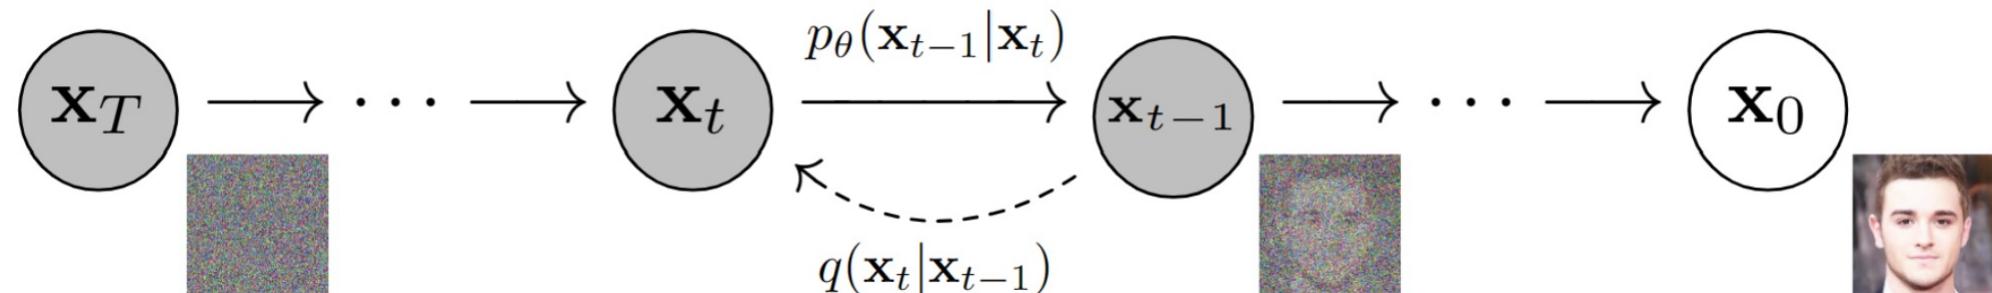
- Training on GT learns to model the data's frequency spectrum

$$\mathcal{L}^k(u, t) = \mathbb{E}_{\epsilon^k \sim \mathcal{N}(0, 1)} [\|\epsilon_k - \text{NO}(u(t) + \sigma_k \epsilon_k, u(t - \Delta t), k)\|_2^2]$$

- Use exponential decreasing noise variance $\sigma_k = \sigma_{\min}^{k/K}$

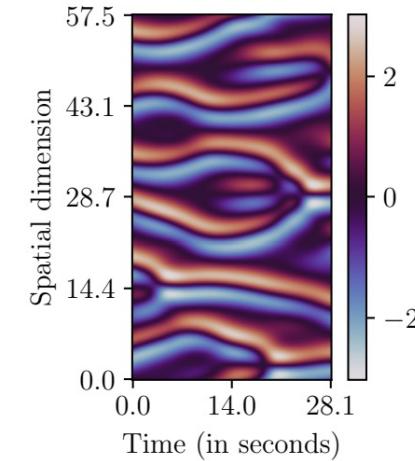
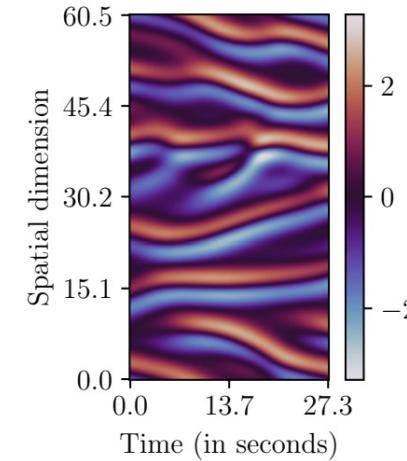
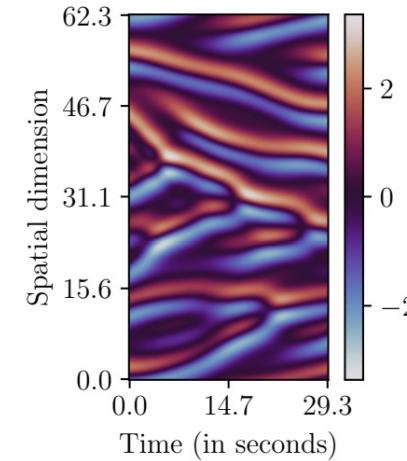
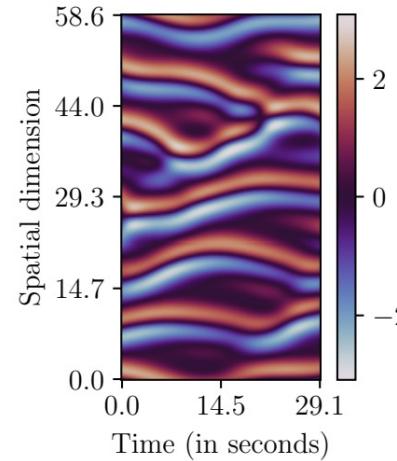
PDE-Refiner – Relation to Diffusion Models

- Popular usage of denoising: Diffusion Models (DDPM) [Ho et al., 2020]
- Key differences of PDE-Refiner to DDPMs:
 1. GT is deterministic \Rightarrow exponential decreasing noise schedule with very small minimum
 2. Speed is of essence for application \Rightarrow very few denoising steps (usually 1-4)
 3. Different objective \Rightarrow predicts signal at initial step



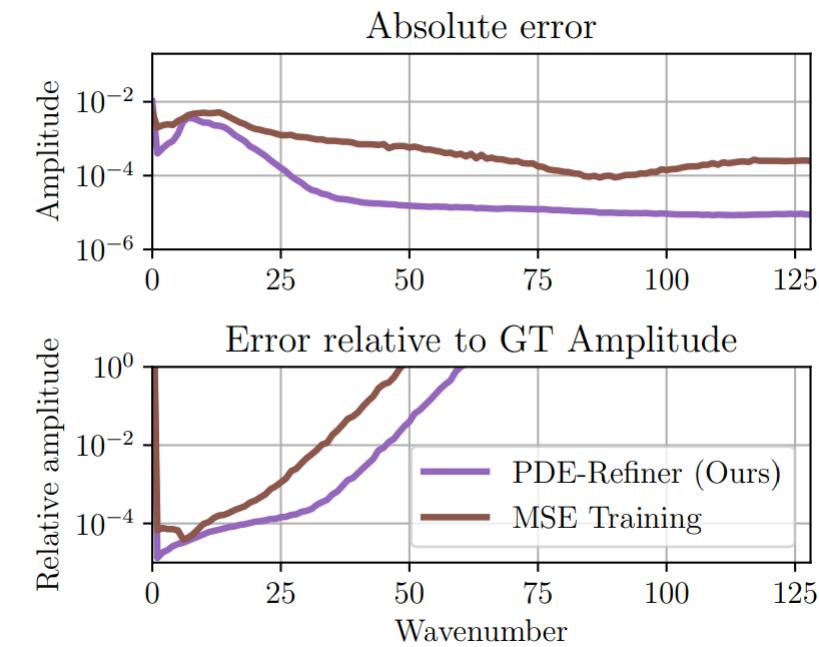
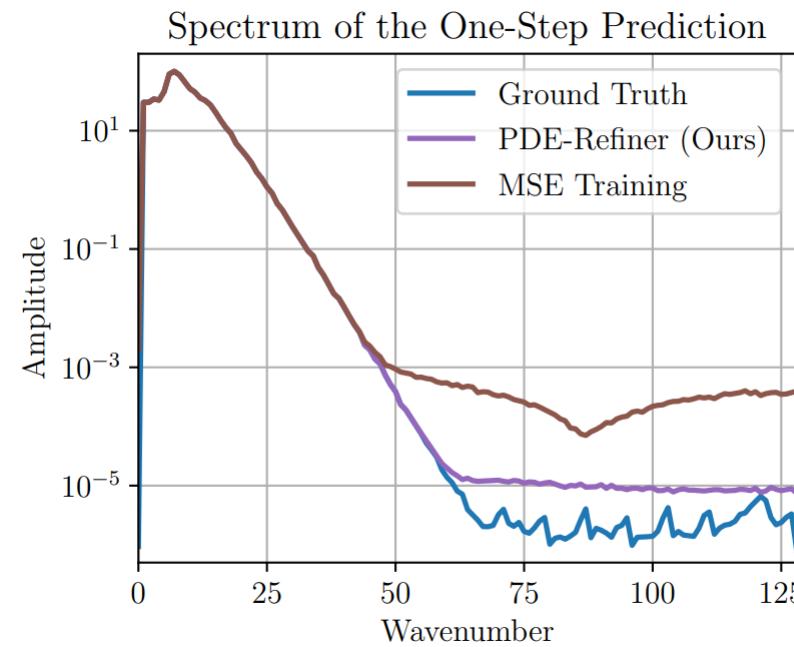
PDE-Refiner – Experimental Setup

- 1D Kuramoto-Sivashinsky Equation
 - Train models on simulated data from classical solver as ground truth
 - Varying initial condition, spatial size, and simulated time step
 - Different neural operator architectures (U-Net, FNO, Dilated ResNets)
 - Evaluation metric: time until correlation between rollout prediction and GT goes below 0.8



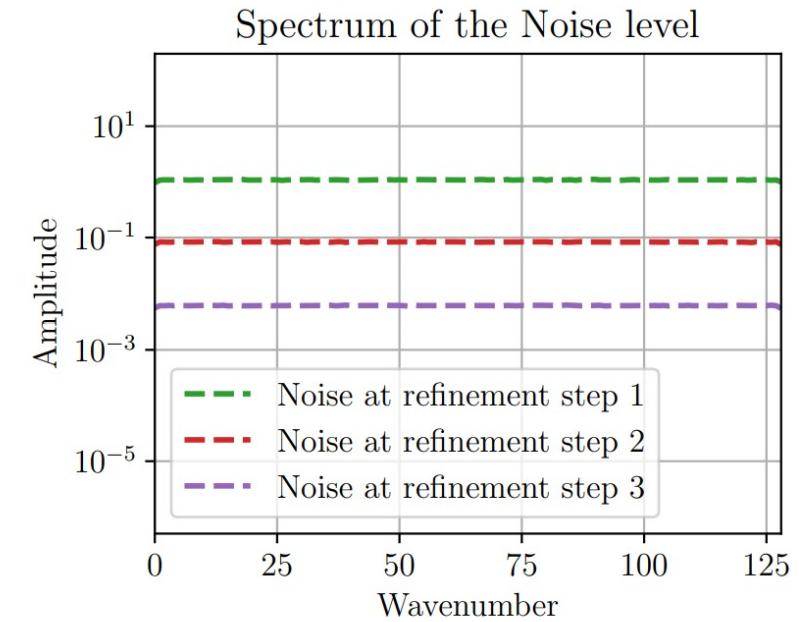
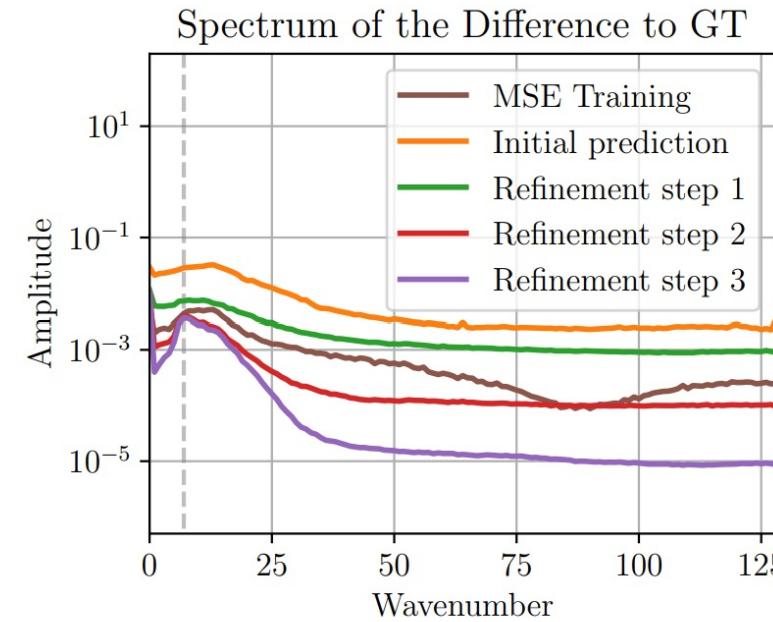
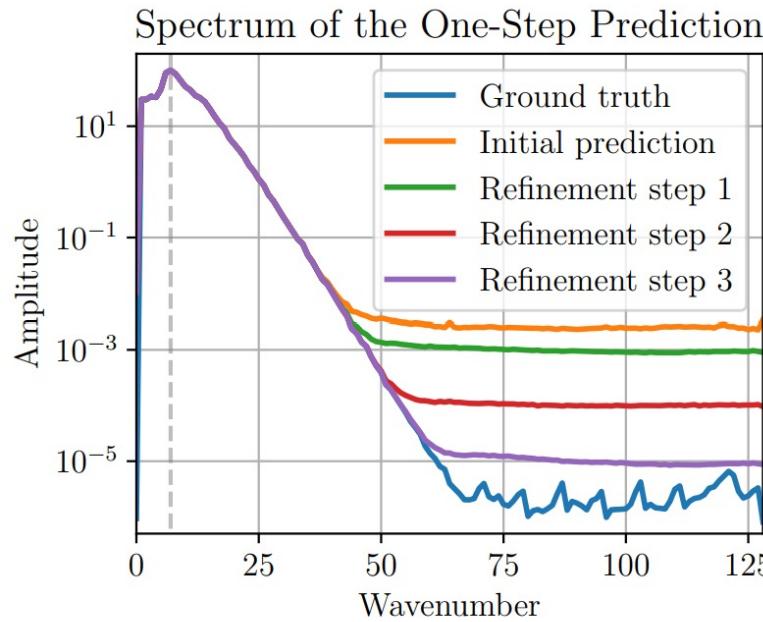
PDE-Refiner – Frequency Spectrum KS equation

- PDE-Refiner models a larger frequency band accurately



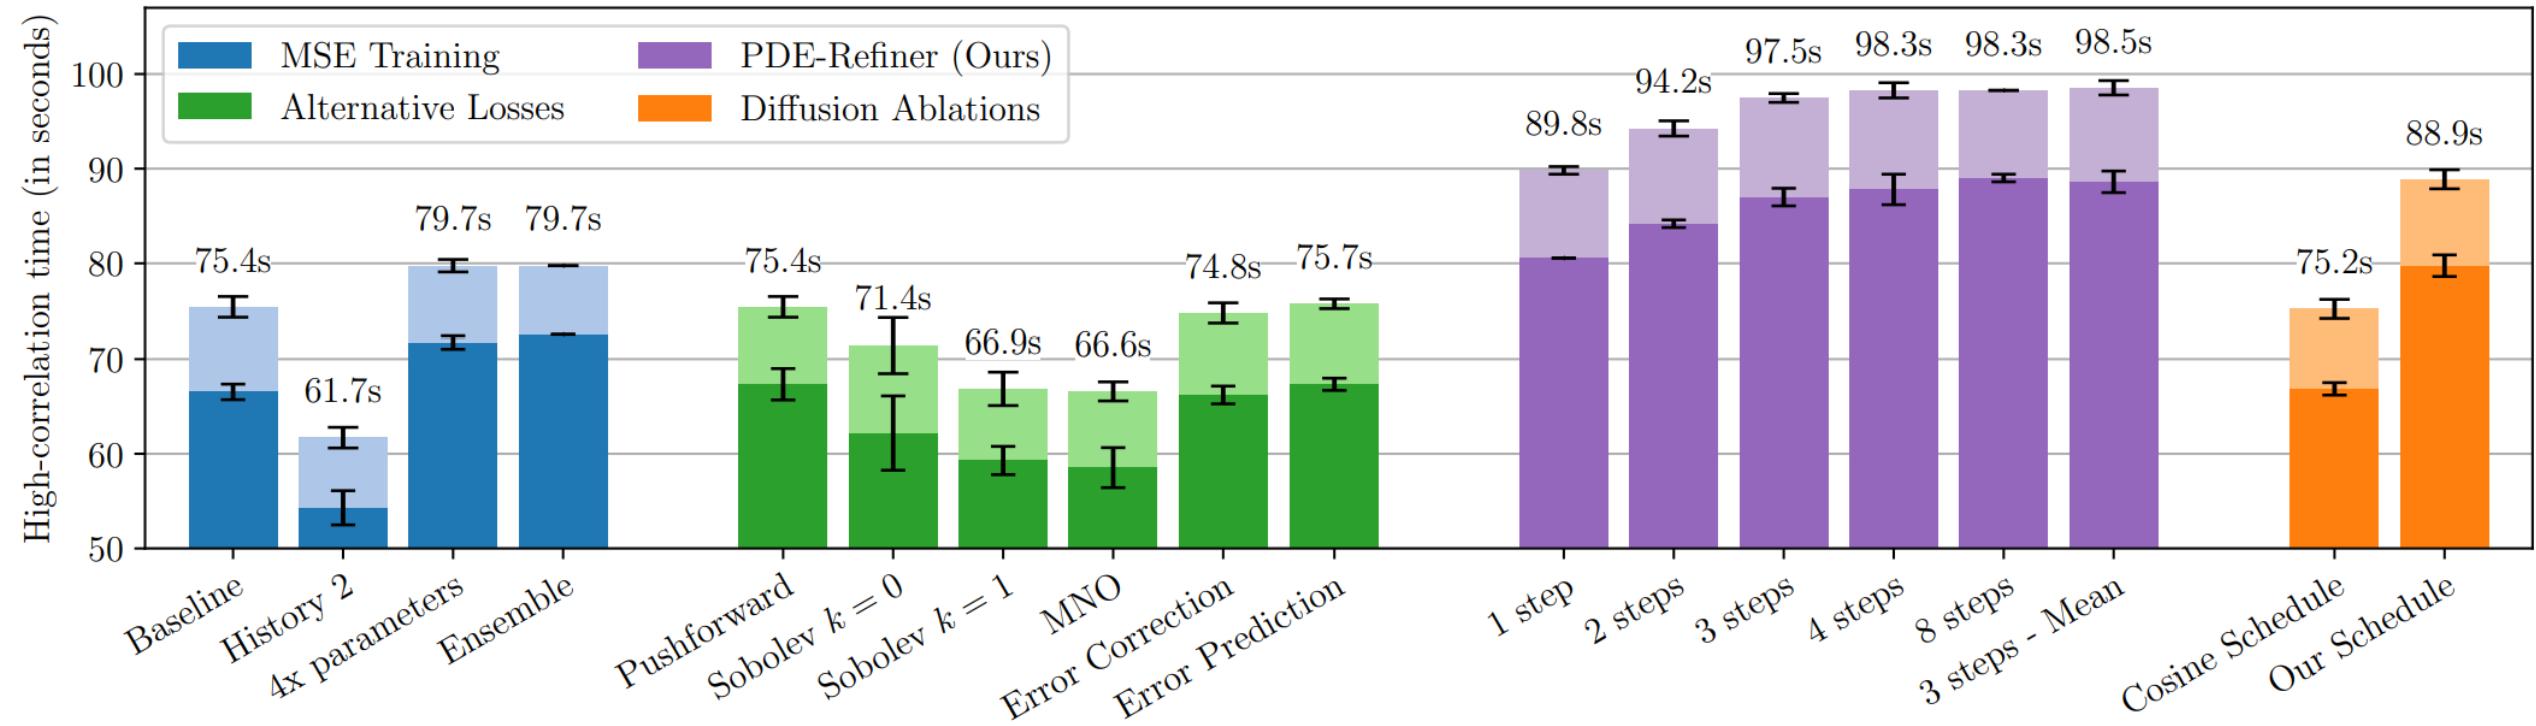
PDE-Refiner – Frequency Spectrum KS equation

- Refinement steps focus on different amplitude levels

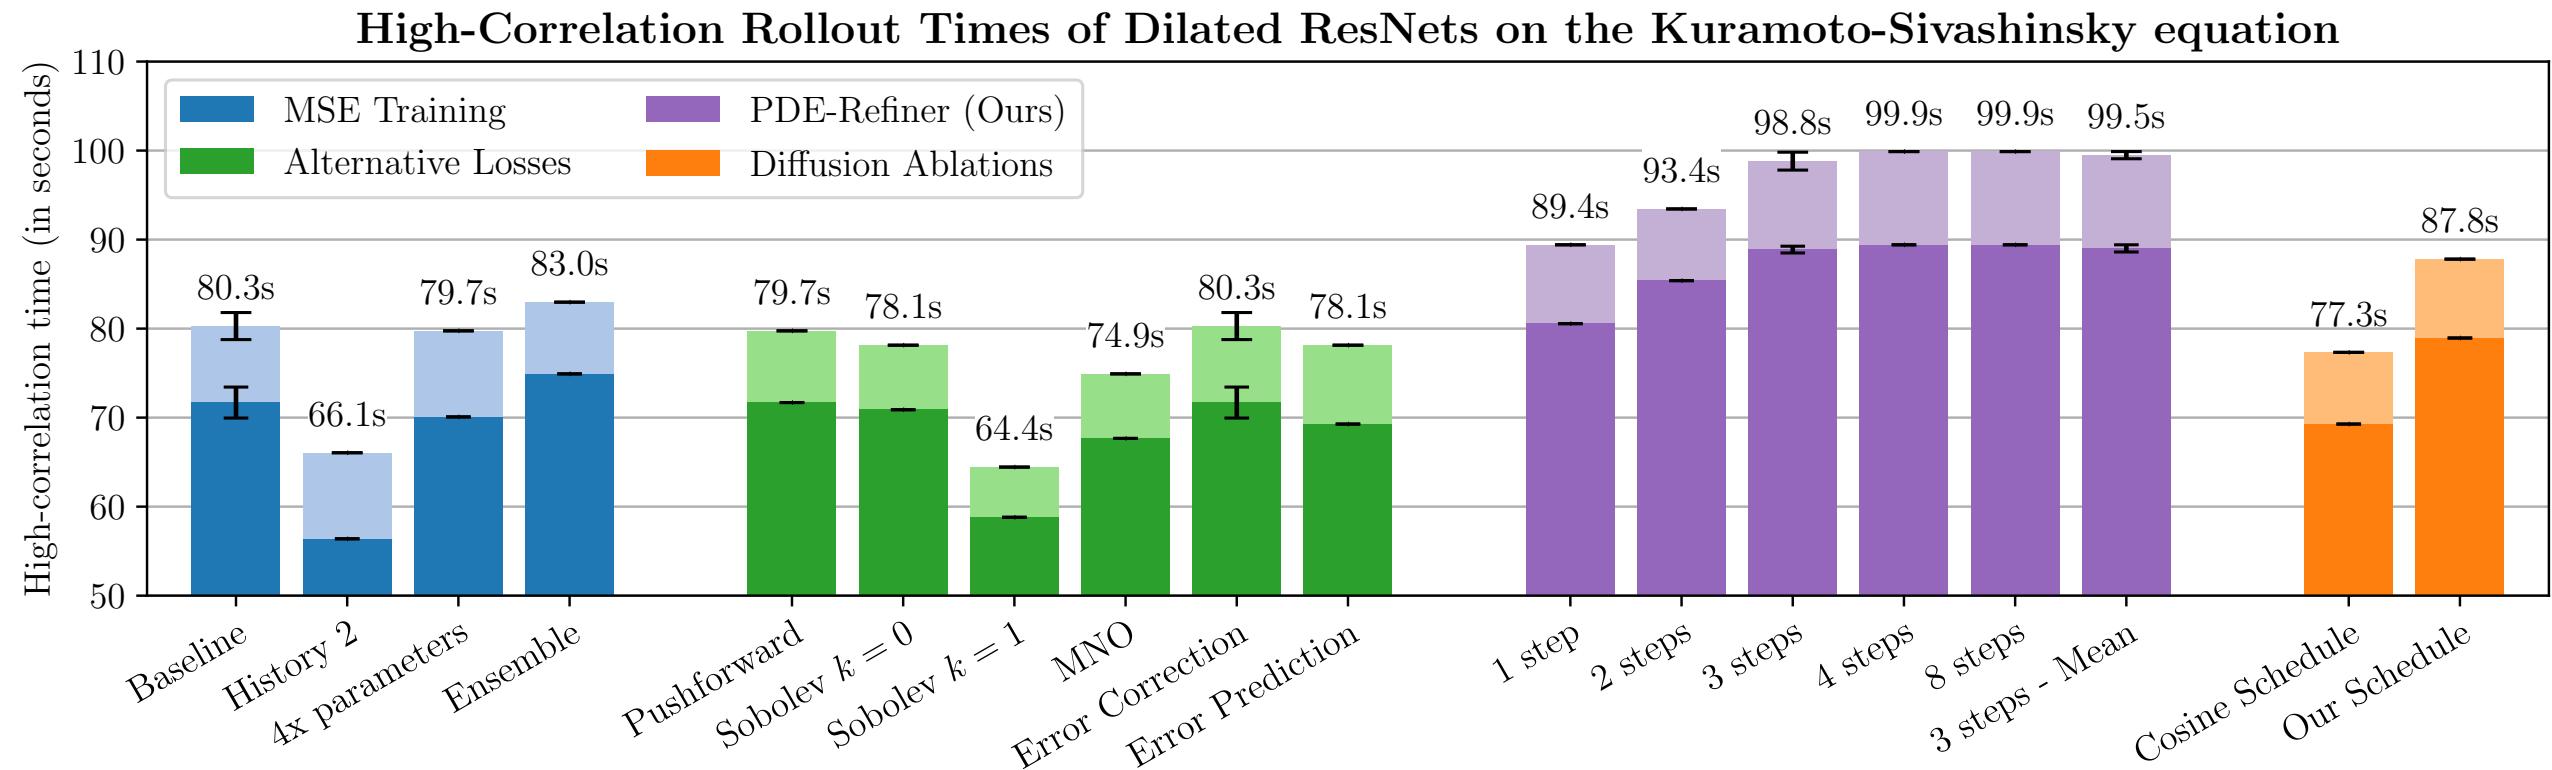


PDE-Refiner – Rollout Performance (U-Net)

High-Correlation Rollout Times on the Kuramoto-Sivashinsky equation

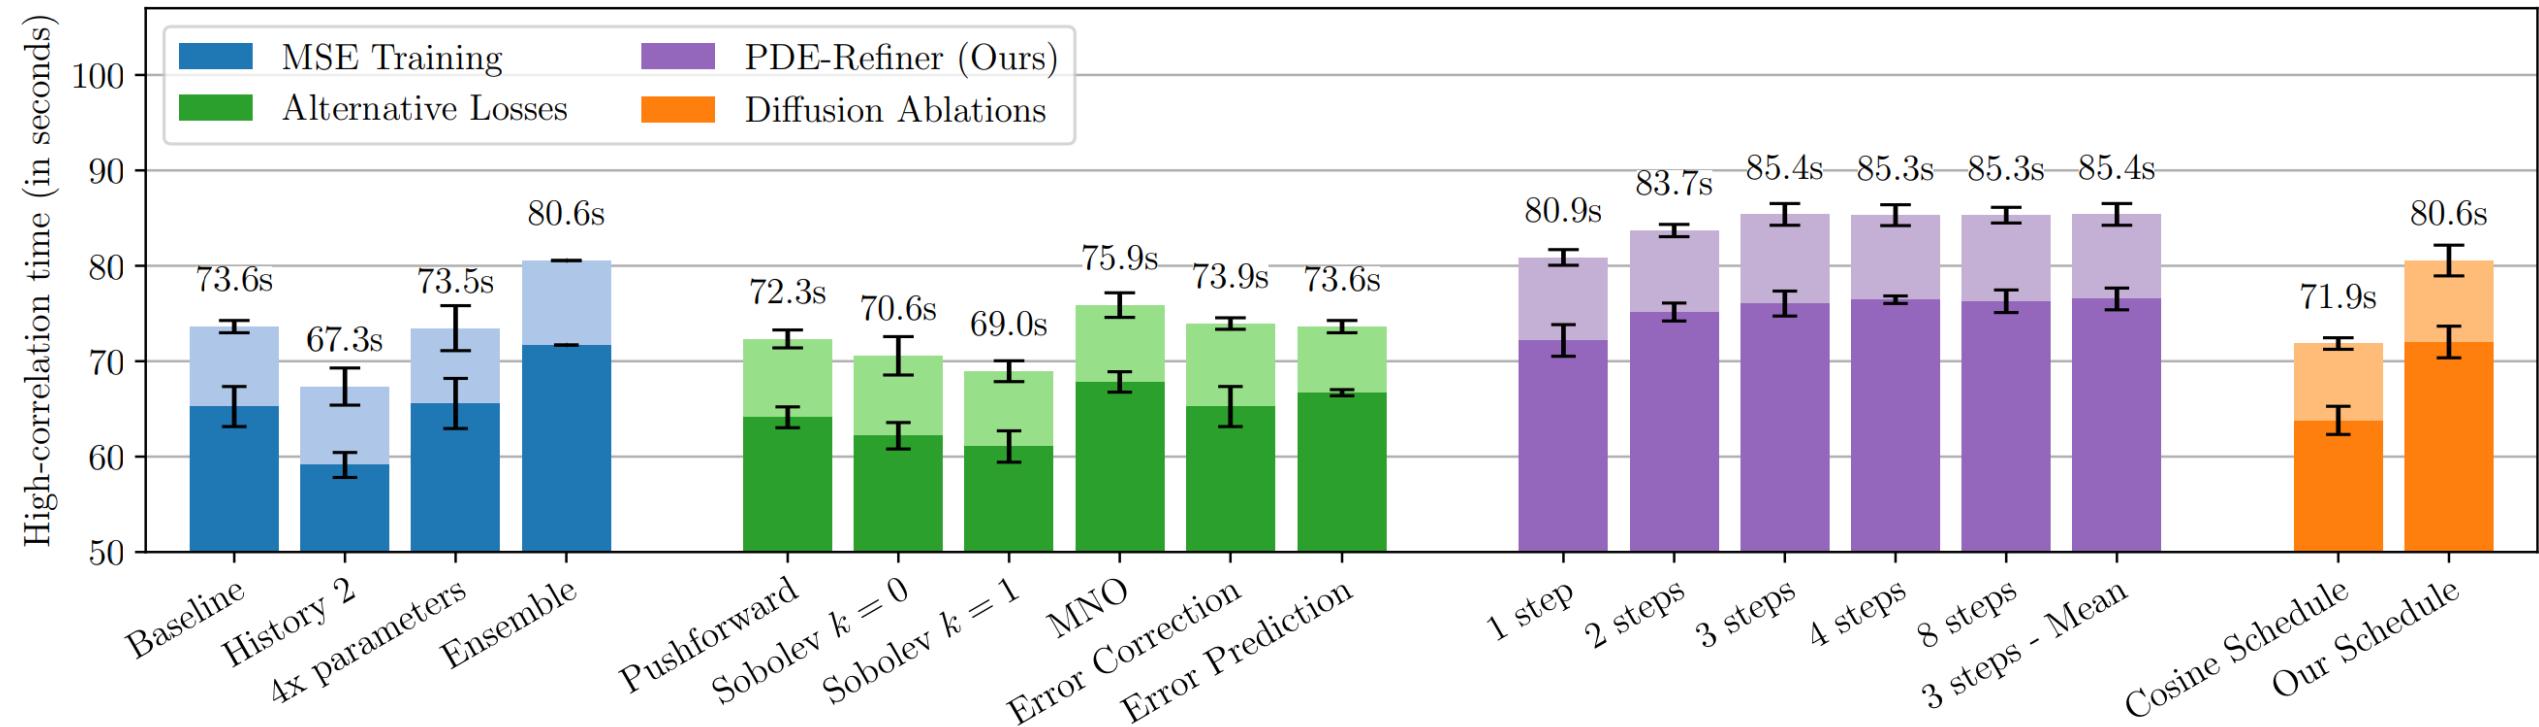


PDE-Refiner – Rollout Performance (Dilated ResNets)



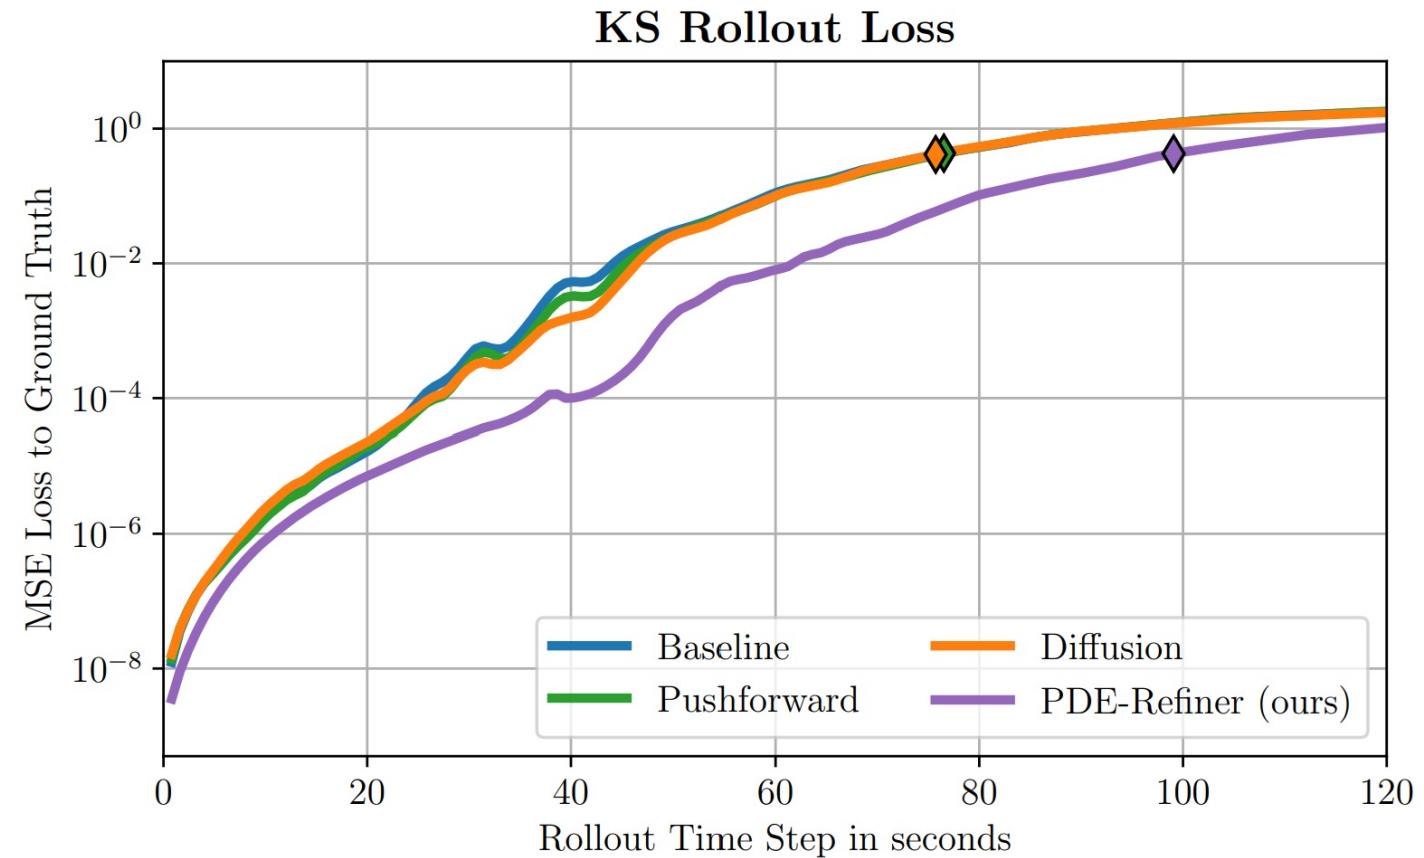
PDE-Refiner – Rollout Performance (FNOs)

High-Correlation Rollout Times of FNOs on the Kuramoto-Sivashinsky equation

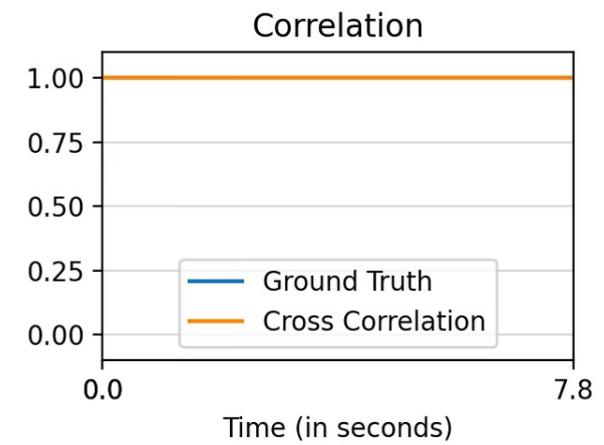
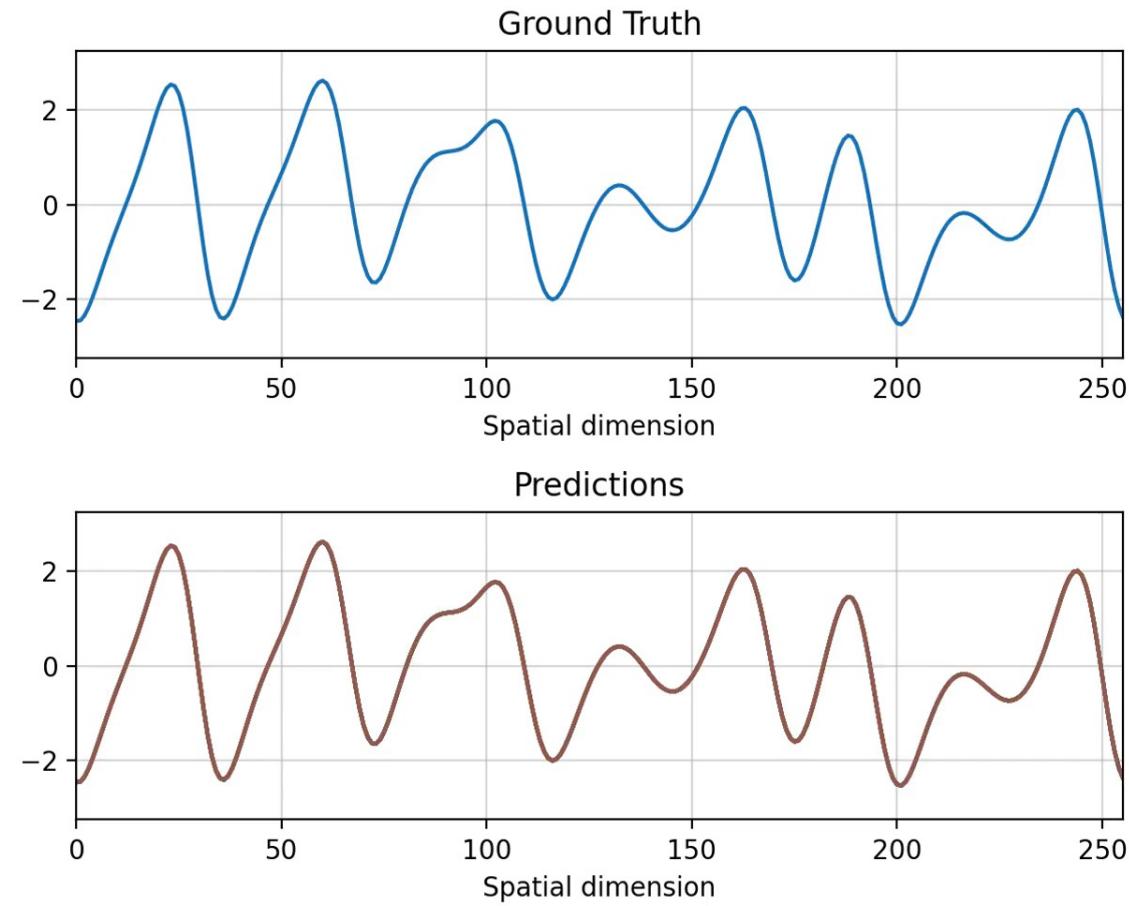


PDE-Refiner – Delaying Error Propagation

- MSE loss of all models similar for first 20 seconds
- Modeling more frequencies allows PDE-Refiner to delay error propagation
- Results in longer stable rollouts

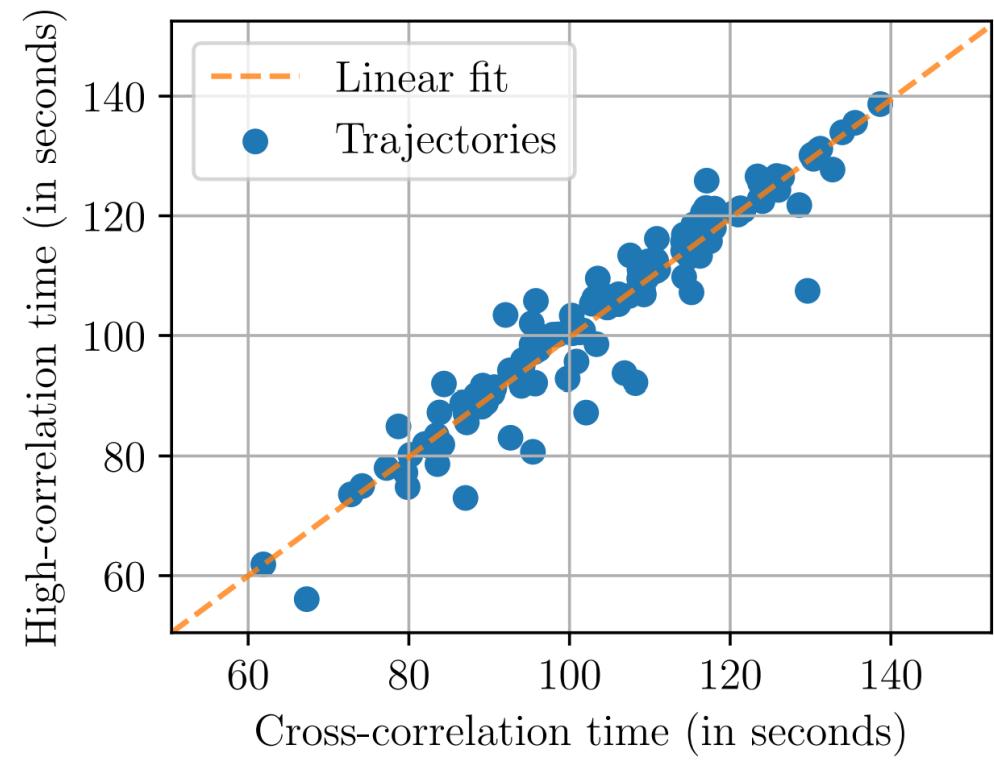


PDE-Refiner – Uncertainty Estimation



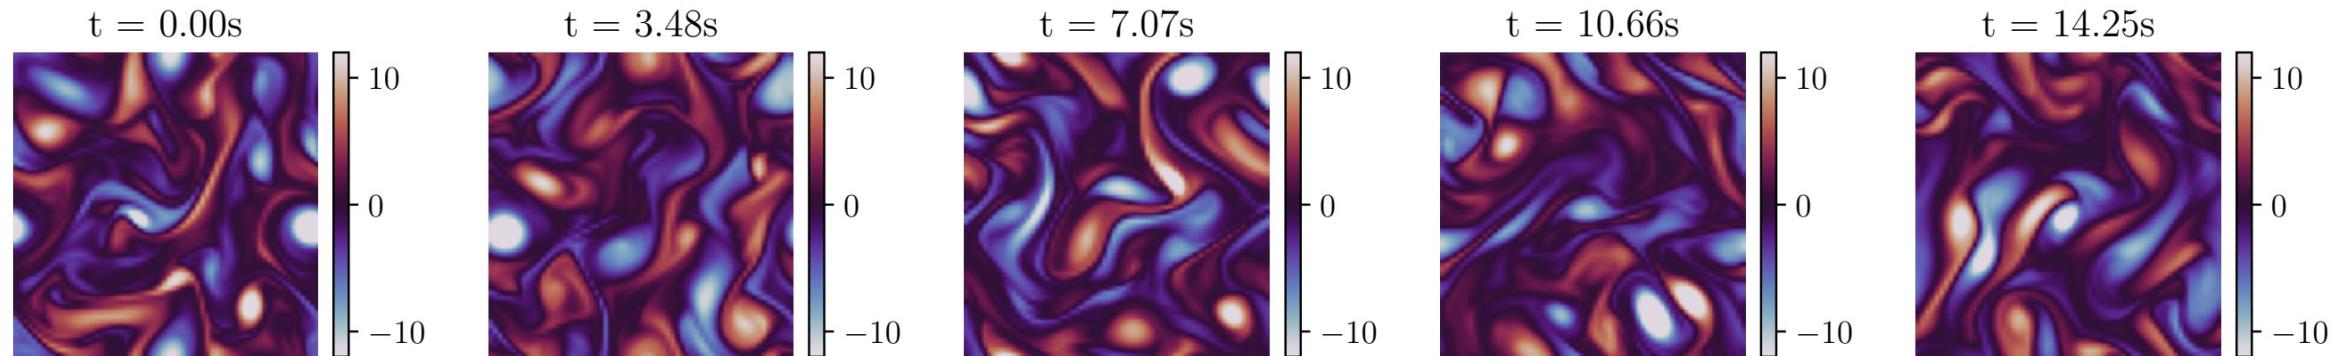
PDE-Refiner – Uncertainty Estimation

- Estimating accurate rollout time by measuring cross-correlation between sampled trajectories
- Accurate uncertainty estimates
- Outperforms other simple uncertainty estimation methods while more efficient than model ensemble



PDE-Refiner – 2D Kolmogorov Flow

- 2D Kolmogorov Flow
 - Generated at 2048x2048 resolution, downscaled to 64x64
 - 2-channel input, evaluated on vorticity



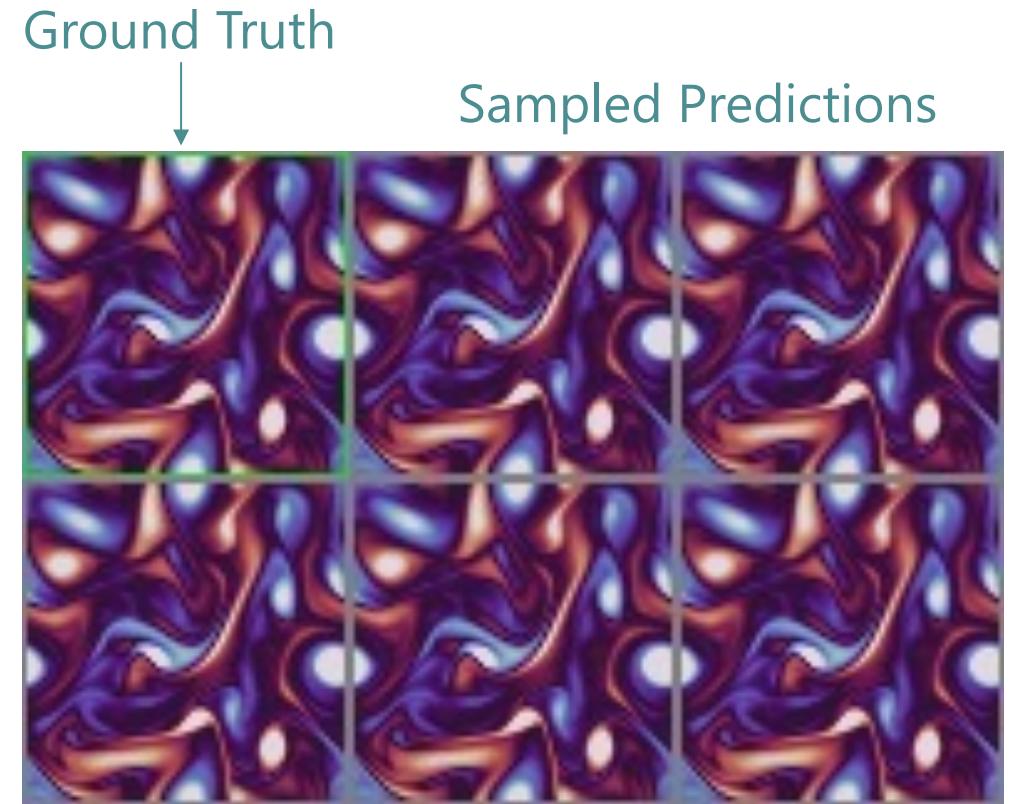
PDE-Refiner – 2D Kolmogorov Flow

- PDE-Refiner outperforms SOTA hybrid solvers
- Gain over MSE baseline smaller than on KS equation
 - Flatter frequency spectrum
 - Smaller resolution
 - Generally higher loss, stronger models possible

Method	Corr. > 0.8 time
<i>Classical PDE Solvers</i>	
DNS - 64×64	2.805
DNS - 128×128	3.983
DNS - 256×256	5.386
DNS - 512×512	6.788
DNS - 1024×1024	8.752
<i>Hybrid Methods</i>	
LC [42, 79] - CNN	6.900
LC [42, 79] - FNO	7.630
LI [42] - CNN	7.910
TSM [75] - FNO	7.798
TSM [75] - CNN	8.359
TSM [75] - HiPPO	9.481
<i>ML Surrogates</i>	
MSE training - FNO	6.451 ± 0.105
MSE training - U-Net	9.663 ± 0.117
PDE-Refiner - U-Net	10.659 ± 0.092

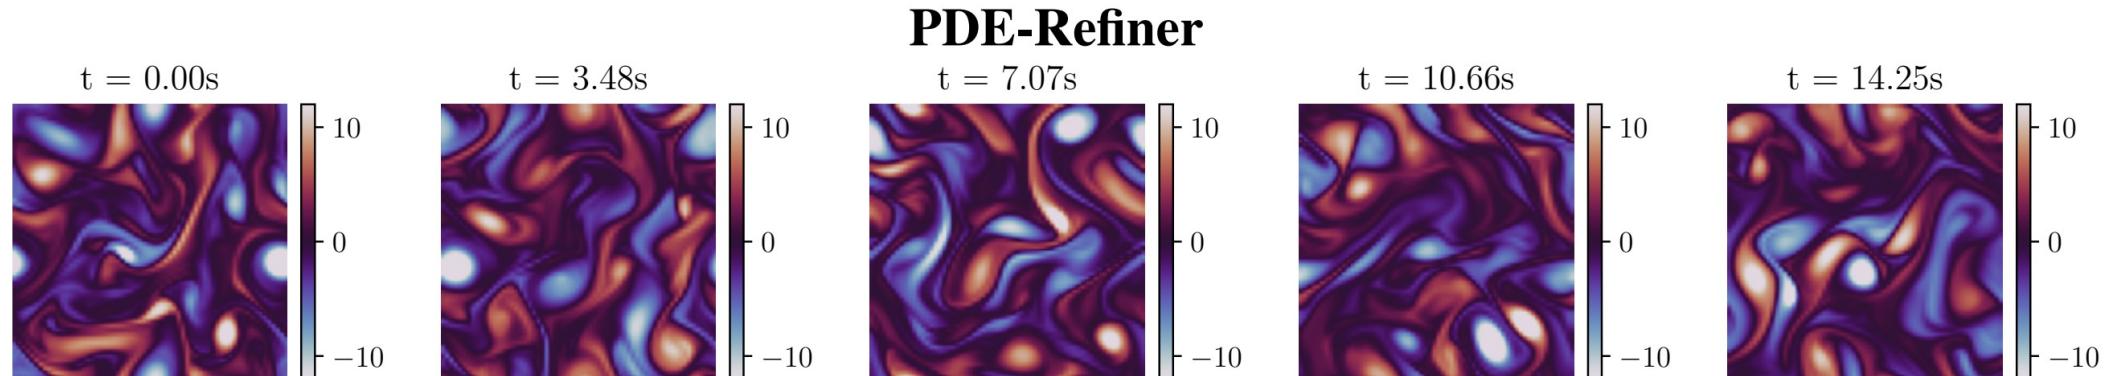
PDE-Refiner – 2D Kolmogorov Flow

- Sampling multiple trajectories estimates uncertainty
- PDE-Refiner slightly overconfident due to small dataset size



PDE-Refiner – Speed Comparison

- Speed comparison on generating 16 trajectories of 20 seconds
 - MSE model: 4 seconds
 - PDE-Refiner: $(K+1) \times \text{MSE} = 16$ seconds
 - Hybrid solver: 20 seconds
 - Classical solver: 31 minutes
- PDE-Refiner offers tunable tradeoff between runtime and accuracy



Summary

- Modeling a large spatial frequency band is key for long accurate rollouts
- PDE-Refiner achieves this by an iterative refinement process, gaining up to 30% longer rollouts across different neural operators and PDEs
- Denoising process inherently learns accurate uncertainty estimate
- PDE-Refiner offers flexible tradeoff between accuracy and speed
- Code coming soon in [PDEArena](#)

PDE-Refiner: Achieving Accurate Long Rollouts with Neural PDE Solvers

Phillip Lippe, Bas S. Veeling, Paris Perdikaris, Richard E. Turner, Johannes Brandstetter

[Project Website](#)

Questions?

